

ADVANCED NETWORK SECURITY AND ARCHITECTURES

METI

Report Module 1 [EN]

Authors:

José Oliveira (103771)

Afonso Mendes (116024)

Tiago Videira (116069)

jose.novais.oliveira@tecnico.ulisboa.pt

afonsofmenDES@tecnico.ulisboa.pt

tiago.g.videira@tecnico.ulisboa.pt

Group 3

Contents

1	Introduction	3
2	Network attacks and countermeasures	4
2.1	Objective	4
2.2	MAC Table Overflow	5
2.2.1	Attack and Countermeasure	5
2.2.2	Switch Comparison	5
2.2.3	Feasibility	6
2.3	DHCP spoofing	6
2.3.1	Attack and Countermeasure	6
2.3.2	AI Prediction	7
2.3.3	Switch Comparison	7
2.3.4	Feasibility	7
2.4	ARP poisoning (MitM)	7
2.4.1	Network Topology	7
2.4.2	Attack and Countermeasure	8
2.4.3	Defense Statistics	9
2.4.4	AI Configuration vs Experimental Setup	9
2.4.5	Switch Comparison	9
2.4.6	Feasibility	10
2.5	DNS spoofing	10
2.6	RIP Poisoning	10
2.6.1	Network Topology	10
2.6.2	Attack and Countermeasure	10
2.6.3	Effectiveness of the Attack	12
2.6.4	AI Prediction vs Experimental Results	13
2.6.5	Feasibility	13
3	Firewalls	14
3.1	Classical firewalls versus Zone-Based Policy Firewalls	14
3.1.1	Firewall Validation	14
3.1.2	Classical Firewall (CBAC)	14
3.1.3	ZBPF	14
3.2	NMAP Operation	15
3.2.1	Host Discovery	15
3.2.2	Port Scanning	15
3.3	CBAC vs ZBPF	15
3.4	Protecting a campus network using a ZBPF	15
3.5	Defense against DoS attacks	15

4	AAA	16
4.1	AAA with TACACS+	16
4.2	AAA with Radius	16
4.3	802.1X Authentication	16
5	Conclusion	17
6	Appendix	18

1 Introduction

FAZER APENAS NO FIM DOS 3 LABS

2 Network attacks and countermeasures

2.1 Objective

The main objective of this laboratory work is to study and experimentally evaluate a set of common network attacks and their corresponding countermeasures in a controlled environment using GNS3.

In particular, the following attacks are addressed:

- **MAC table overflow:** to understand how a switch can be forced to behave as a hub by saturating its MAC address table, and to evaluate mitigation mechanisms such as port security.
- **DHCP spoofing:** to demonstrate how a rogue DHCP server can provide malicious network configurations to clients, and to analyze the effectiveness of DHCP snooping as a countermeasure.
- **ARP poisoning (Man-in-the-Middle):** to perform a MitM attack by manipulating ARP tables, and to study the use of Dynamic ARP Inspection (DAI) for protection.
- **DNS spoofing:** to analyze how a victim can be redirected to a fake web server through DNS manipulation combined with ARP poisoning.
- **RIP poisoning:** to investigate how routing tables can be corrupted using forged RIP messages, leading to traffic redirection, and to explore authentication-based defenses.

Additionally, this work aims to:

- Analyze network behavior through packet captures using Wireshark;
- Interpret device-level information using Cisco IOS commands;
- Evaluate the practical feasibility and limitations of each attack;
- Compare experimental results with theoretical expectations.

2.2 MAC Table Overflow


2.2.1 Attack and Countermeasure

The MAC table overflow attack was performed using the `camov.py` script, generating a large number of frames with different spoofed MAC addresses.

As observed in Wireshark (Figure 1), a high volume of traffic with varying source addresses was captured, indicating the flooding behavior of the attack.

The switch behavior was analyzed using `show mac address-table`.

To mitigate the attack, port security was configured, limiting the number of MAC addresses per port and preventing the insertion of multiple spoofed entries.



No.	Time	Source	Destination	Protocol	Length	Info
142	13.578719	64.4.149.221	218.228.255.127	IPv4	60	
143	13.575639	95.116.174.88	250.1.202.238	IPv4	60	
144	13.607635	93.2.163.132	118.119.242.91	IPv4	60	
145	13.623101	190.38.44.3	142.186.181.76	IPv4	60	
146	13.672099	79.241.71.185	128.238.207.11	IPv4	60	
147	13.679491	108.62.180.42	198.23.220.199	IPv4	60	
148	13.716949	235.42.210.44	197.222.120.239	IPv4	60	
149	13.721042	170.104.88.50	199.159.79.108	IPv4	60	
150	13.732057	33.0.137.130	136.43.59.211	IPv4	60	
151	13.735855	145.227.196.138	198.63.73.5	IPv4	60	
152	13.765794	7.137.177.198	251.206.70.219	IPv4	60	
153	13.769912	213.182.115.152	235.44.132.143	IPv4	60	
154	13.798114	130.143.108.247	216.253.46.70	IPv4	60	
155	13.802038	68.137.63.172	240.13.34.126	IPv4	60	
156	13.828729	59.132.235.29	59.134.240.32	IPv4	60	
157	13.836100	65.48.16.31	104.66.43.26	IPv4	60	
158	13.865297	113.128.218.153	64.89.86.3	IPv4	60	
159	13.869642	228.226.232.14	183.205.234.29	IPv4	60	
160	13.898473	151.215.76.250	52.12.174.145	IPv4	60	
161	13.902513	164.165.101.160	240.206.230.245	IPv4	60	
162	13.930701	194.108.216.110	95.241.143.222	IPv4	60	
163	13.934569	55.228.93.71	36.163.107.232	IPv4	60	
164	13.962949	42.218.8.114	132.71.23.74	IPv4	60	
165	13.967965	17.104.180.209	52.75.122.66	IPv4	60	
166	13.996842	61.30.141.106	93.5.65.70	IPv4	60	
167	14.000613	99.81.18.179	62.20.13.194	IPv4	60	
168	14.014706	83.205.144.209	231.16.1.98	IPv4	60	
169	14.018855	62.109.206.157	226.80.27.63	IPv4	60	
171	14.030336	218.217.66.104	235.118.255.196	IPv4	60	
172	14.034170	21.56.250.252	128.194.52.253	IPv4	60	
173	14.078302	29.87.171.46	109.67.239.132	IPv4	60	
174	14.081964	102.166.190.135	57.149.151.175	IPv4	60	

Figure 1: Wireshark capture during MAC table overflow attack

2.2.2 Switch Comparison

Cisco Catalyst 2960 supports port security, allowing the configuration of a maximum number of MAC addresses per port and defining actions when violations occur (e.g., restrict or shutdown). This makes it possible to prevent MAC table overflow attacks.

Link: www.cisco.com/c/en/us/products/switches/catalyst-2960-series-switches/index.html

In contrast, the TP-Link TL-SF1005D is an unmanaged switch that does not support port security or any traffic control mechanisms, as it lacks a configuration interface. Therefore, it cannot prevent MAC flooding attacks.

Link: www.tp-link.com/en/home-networking/soho-switch/tl-sf1005d/

2.2.3 Feasibility

The objective of this attack is to flood the switch MAC address table, forcing the switch to behave like a hub. In that state, frames may be forwarded to all ports, allowing the attacker to listen to conversations between other hosts.

However, in practice, the attack may be difficult to execute successfully on modern switches because their MAC address tables are large. Additionally, the high amount of generated traffic can cause network instability or a denial-of-service condition, instead of allowing the attacker to reliably capture useful traffic.

Therefore, although the attack is theoretically possible, its practical feasibility is limited and depends on the switch capacity, configuration, and security features.

2.3 DHCP spoofing

2.3.1 Attack and Countermeasure


The DHCP spoofing attack was performed by introducing a rogue DHCP server that responds to DHCP Discover messages with malicious configuration parameters.

As observed in Wireshark (Figure 2), multiple DHCP Offer messages were received for the same transaction ID, indicating a race condition between the legitimate and the rogue DHCP servers. The victim accepts the first Offer received, which may result in a malicious gateway configuration provided by the attacker.

20	20.480048	20.0.0.100	20.0.0.200	DHCP	338 DHCP Offer	- Transaction ID 0x2c5ff044
21	20.482689	20.0.0.2	20.0.0.1	DHCP	342 DHCP Offer	- Transaction ID 0x2c5ff044
23	21.447930	255.255.255.255	0.0.0.0	DHCP	406 DHCP Request	- Transaction ID 0x2c5ff044
24	21.493701	20.0.0.100	20.0.0.200	DHCP	338 DHCP ACK	- Transaction ID 0x2c5ff044

Figure 2: Wireshark capture showing multiple DHCP Offer messages (attack)

After enabling DHCP snooping and configuring the port connected to the legitimate DHCP server as trusted, only one DHCP Offer was observed (Figure 3). The rogue DHCP responses were blocked by the switch, and the client received a valid configuration. This confirms that DHCP snooping effectively mitigates the attack.



The image shows a Wireshark capture window titled '*Standard input [Switch1 Ethernet2 to CiscoIOSvL215.2(4.0.55)E-1 Gi0/3]'. The packet list shows five DHCP packets. The packet details pane shows the selected packet (No. 7) as a DHCP Offer with Transaction ID 0xc74c155c. The packet bytes pane shows the DHCP Offer structure.

No.	Time	Destination	Source	Protocol	Length	Info
6	8.697745	255.255.255.255	0.0.0.0	DHCP	406	DHCP Discover - Transaction ID 0xc74c155c
7	8.714163	20.0.0.2	20.0.0.1	DHCP	342	DHCP Offer - Transaction ID 0xc74c155c
8	9.676498	255.255.255.255	0.0.0.0	DHCP	406	DHCP Request - Transaction ID 0xc74c155c
9	9.685474	20.0.0.2	20.0.0.1	DHCP	342	DHCP ACK - Transaction ID 0xc74c155c
12	11.279396	255.255.255.255	0.0.0.0	DHCP	406	DHCP Discover - Transaction ID 0x9d2bb165

Figure 3: Wireshark capture with DHCP snooping enabled (only legitimate Offer)

2.3.2 AI Prediction

The AI predicted that the victim closer to the attacker would be more likely to be compromised, since the rogue DHCP Offer would arrive first. This prediction matched the experimental results, where the victim receiving the first Offer obtained the malicious configuration.

2.3.3 Switch Comparison

Cisco Catalyst 2960 supports DHCP snooping, allowing ports to be classified as trusted or untrusted and blocking DHCP server messages from untrusted ports.

Link: www.cisco.com/c/en/us/products/switches/catalyst-2960-series-switches/index.html

The TP-Link TL-SF1005D does not support DHCP snooping, as it is an unmanaged switch without configuration capabilities or security features.

Link: www.tp-link.com/en/home-networking/soho-switch/tl-sf1005d/

2.3.4 Feasibility

The attack is feasible in networks without DHCP snooping, especially when the attacker is closer to the victim than the legitimate DHCP server. However, due to the race condition, success is not guaranteed. When DHCP snooping is enabled, the attack becomes ineffective.

2.4 ARP poisoning (MitM)

2.4.1 Network Topology

The ARP poisoning attack was performed in the network shown in Figure 4, where two hosts communicate through a switch and an attacker is connected to the same Layer 2 domain.




Figure 4: Network topology for ARP poisoning attack

2.4.2 Attack and Countermeasure

As observed in Wireshark (Figure 5), ARP packets indicate duplicate IP usage, confirming that the attacker is associating its MAC address with multiple IP addresses.

At the same time, ICMP traffic between the victims is still observed, which demonstrates that communication is maintained while being intercepted by the attacker, confirming a successful Man-in-the-Middle (MitM) condition.




Figure 5: ARP poisoning showing duplicate IP detection and intercepted ICMP traffic

The attacker generated multiple forged ARP requests and replies, attempting to poison the ARP tables of the victims, as shown in Figure 6.

22	41.442250	Broadcast	02:42:e1:18:9f:00	ARP	42	Who has 20.0.0.3? Tell 20.0.0.2
24	43.473802	Broadcast	02:42:e1:18:9f:00	ARP	42	Who has 20.0.0.4? Tell 20.0.0.2
26	45.523090	Broadcast	02:42:e1:18:9f:00	ARP	42	Who has 20.0.0.3? Tell 20.0.0.2
28	47.548118	Broadcast	02:42:e1:18:9f:00	ARP	42	20.0.0.4 is at 02:42:e1:18:9f:00
29	47.597578	Broadcast	02:42:e1:18:9f:00	ARP	42	Who has 20.0.0.4? Tell 20.0.0.2
31	49.613918	Broadcast	02:42:e1:18:9f:00	ARP	42	20.0.0.3 is at 02:42:e1:18:9f:00
32	51.161927	Broadcast	02:42:e1:18:9f:00	ARP	42	Who has 20.0.0.3? Tell 20.0.0.2
34	53.178241	Broadcast	02:42:e1:18:9f:00	ARP	42	20.0.0.4 is at 02:42:e1:18:9f:00
35	53.226972	Broadcast	02:42:e1:18:9f:00	ARP	42	Who has 20.0.0.4? Tell 20.0.0.2
37	55.245424	Broadcast	02:42:e1:18:9f:00	ARP	42	20.0.0.3 is at 02:42:e1:18:9f:00
41	56.794383	Broadcast	02:42:e1:18:9f:00	ARP	42	Who has 20.0.0.3? Tell 20.0.0.2
43	58.809905	Broadcast	02:42:e1:18:9f:00	ARP	42	20.0.0.4 is at 02:42:e1:18:9f:00
44	58.862239	Broadcast	02:42:e1:18:9f:00	ARP	42	Who has 20.0.0.4? Tell 20.0.0.2

Figure 6: ARP packets generated by the attacker

However, the attack was mitigated using Dynamic ARP Inspection (DAI). The switch detected invalid ARP packets and blocked them, as shown in the system logs (Figure 7).

```
*Apr 29 10:34:18.724: %SW_DAI-4-DHCP_SNOOPING_DENY: 1 Invalid ARPs (Req) on Gi0/0, vlan 1.([0242.e118.9f00/20.0.0.2/0000.000
0.0000/20.0.0.3/10:34:18 UTC Wed Apr 29 2026])
*Apr 29 10:34:20.876: %SW_DAI-4-DHCP_SNOOPING_DENY: 1 Invalid ARPs (Res) on Gi0/0, vlan 1.([0242.e118.9f00/20.0.0.4/0000.000
0.0000/20.0.0.3/10:34:20 UTC Wed Apr 29 2026])
*Apr 29 10:34:20.877: %SW_DAI-4-DHCP_SNOOPING_DENY: 1 Invalid ARPs (Req) on Gi0/0, vlan 1.([0242.e118.9f00/20.0.0.2/0000.000
0.0000/20.0.0.4/10:34:20 UTC Wed Apr 29 2026])
*Apr 29 10:34:23.095: %SW_DAI-4-DHCP_SNOOPING_DENY: 1 Invalid ARPs (Res) on Gi0/0, vlan 1.([0242.e118.9f00/20.0.0.3/0000.000
0.0000/20.0.0.4/10:34:21 UTC Wed Apr 29 2026])
*Apr 29 10:34:24.097: %SW_DAI-4-DHCP_SNOOPING_DENY: 1 Invalid ARPs (Req) on Gi0/0, vlan 1.([0242.e118.9f00/20.0.0.2/0000.000
0.0000/20.0.0.3/10:34:23 UTC Wed Apr 29 2026])
*Apr 29 10:34:25.279: %SW_DAI-4-DHCP_SNOOPING_DENY: 1 Invalid ARPs (Res) on Gi0/0, vlan 1.([0242.e118.9f00/20.0.0.4/0000.000
0.0000/20.0.0.3/10:34:25 UTC Wed Apr 29 2026])
*Apr 29 10:34:25.280: %SW_DAI-4-DHCP_SNOOPING_DENY: 1 Invalid ARPs (Req) on Gi0/0, vlan 1.([0242.e118.9f00/20.0.0.2/0000.000
0.0000/20.0.0.4/10:34:25 UTC Wed Apr 29 2026])
*Apr 29 10:34:26.622: %CDP-4-DUPLEX_MISMATCH: duplex mismatch discovered on GigabitEthernet0/3 (not half duplex), with DHCP
FastEthernet0/0 (half duplex).
```

Figure 7: DAI log showing invalid ARP packets being denied

2.4.3 Defense Statistics

The effectiveness of DAI is confirmed by the statistics shown in Figure 8, where a high number of ARP packets were dropped by the switch.

```
Switch#show ip arp inspection statistics
```

Vlan	Forwarded	Dropped	DHCP Drops	ACL Drops
---	-----	-----	-----	-----
1	6	259	259	0

Vlan	DHCP Permits	ACL Permits	Probe Permits	Source MAC Failures
---	-----	-----	-----	-----
1	6	0	0	0

Vlan	Dest MAC Failures	IP Validation Failures	Invalid Protocol Data
---	-----	-----	-----
1	0	0	0

Figure 8: DAI statistics showing dropped ARP packets

2.4.4 AI Configuration vs Experimental Setup

The AI suggested enabling DHCP snooping and configuring Dynamic ARP Inspection on the switch, marking trusted interfaces accordingly.

This matched the experimental setup, where DHCP snooping was required for DAI to validate ARP packets. The configuration had to be adapted to the topology, particularly in defining trusted ports and VLAN configuration.

2.4.5 Switch Comparison

Cisco Catalyst 2960 supports Dynamic ARP Inspection, allowing validation of ARP packets based on DHCP snooping bindings.

Link: www.cisco.com/c/en/us/products/switches/catalyst-2960-series-switches/index.html

The TP-Link TL-SF1005D does not support DAI, as it is an unmanaged switch without security or inspection mechanisms.

Link: www.tp-link.com/en/home-networking/soho-switch/tl-sf1005d/

2.4.6 Feasibility

The attacker attempted to poison the ARP tables by sending forged ARP messages. The attack successfully created a Man-in-the-Middle condition, as observed in Wireshark. However, based on the switch logs and DAI statistics, these packets were detected and blocked.

Therefore, although the attack is feasible in networks without protection, in this setup it was not successful due to the correct configuration of Dynamic ARP Inspection.

2.5 DNS spoofing

2.6 RIP Poisoning

2.6.1 Network Topology

The RIP poisoning attack was performed using the topology shown in Figure 9. The Web browser, the legitimate Web server and the fake Web server are located in different subnets, with the attacker collocated with the fake Web server.




Figure 9: Network topology used for the RIP poisoning attack

2.6.2 Attack and Countermeasure

The attacker sent forged RIPv2 Response messages in order to poison the routing tables of the routers. These messages advertised a route to the legitimate Web server network through the attacker, causing traffic from the browser to be redirected to the fake Web server.

As shown in Figure 10, Wireshark captured multiple RIPv2 Response packets sent to the multicast address 224.0.0.9, which is used by RIPv2 routers. The captured packet advertises the network 10.0.0.0 with metric 5 and subnet mask 255.255.255.128, showing that the attacker injected routing information into the network.

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000	c2:02:37:fd:00:10	c2:02:37:fd:00:10	LOOP	60	Reply
2	1.733278	172.16.0.2	224.0.0.9	RIPv2	66	Response
3	3.737119	172.16.0.2	224.0.0.9	RIPv2	66	Response
4	5.740474	172.16.0.2	224.0.0.9	RIPv2	66	Response
5	6.658340	192.168.1.2	10.0.0.2	TCP	74	54068 → 80 [SYN] Seq=0 Win=6424
6	6.658522	10.0.0.2	192.168.1.2	TCP	74	80 → 54068 [SYN, ACK] Seq=0 Ack=
7	6.688746	192.168.1.2	10.0.0.2	TCP	66	54068 → 80 [ACK] Seq=1 Ack=1 Wi
8	6.698864	192.168.1.2	10.0.0.2	HTTP	476	GET / HTTP/1.1
9	6.699201	10.0.0.2	192.168.1.2	TCP	66	80 → 54068 [ACK] Seq=1 Ack=411
10	6.699222	10.0.0.2	192.168.1.2	HTTP	254	HTTP/1.1 304 Not Modified
11	6.739220	192.168.1.2	10.0.0.2	TCP	66	54068 → 80 [ACK] Seq=411 Ack=18
12	7.744338	172.16.0.2	224.0.0.9	RIPv2	66	Response
13	9.748365	172.16.0.2	224.0.0.9	RIPv2	66	Response
14	9.991713	c2:02:37:fd:00:10	c2:02:37:fd:00:10	LOOP	60	Reply
15	11.746292	02:42:b2:9b:6c:00	c2:02:37:fd:00:10	ARP	42	Who has 172.16.0.1? Tell 172.16
16	11.751331	172.16.0.2	224.0.0.9	RIPv2	66	Response
17	11.752374	c2:02:37:fd:00:10	02:42:b2:9b:6c:00	ARP	60	172.16.0.1 is at c2:02:37:fd:00
18	13.755114	172.16.0.2	224.0.0.9	RIPv2	66	Response
19	15.758195	172.16.0.2	224.0.0.9	RIPv2	66	Response
20	15.898924	172.16.0.1	224.0.0.9	RIPv2	126	Response
21	16.875705	192.168.1.2	10.0.0.2	TCP	66	[TCP Keep-Alive] 54068 → 80 [AC
22	16.875934	10.0.0.2	192.168.1.2	TCP	66	[TCP Keep-Alive ACK] 80 → 54068
23	17.760482	172.16.0.2	224.0.0.9	RIPv2	66	Response
24	17.762445	172.16.0.2	224.0.0.9	RIPv2	66	Response

Source Port: 520
Destination Port: 520
Length: 32
Checksum: 0x63f8 [unverified]
[Checksum Status: Unverified]
[Stream index: 0]
▸ [Timestamps]
UDP payload (24 bytes)
▾ Routing Information Protocol
Command: Response (2)
Version: RIPv2 (2)
▾ IP Address: 10.0.0.0, Metric: 5
Address Family: IP (2)
Route Tag: 0
IP Address: 10.0.0.0
Netmask: 255.255.255.128
Next Hop: 0.0.0.0
Metric: 5

Figure 10: Forged RIPv2 response advertising a route to the target network

The effect of the attack was confirmed in the browser. As shown in Figure 11, the Web browser was redirected to the fake Web server, displaying the message “FAKE SERVER - ATTACK SUCCESS”.




Figure 11: Browser redirected to the fake Web server

Additional Wireshark captures show TCP and HTTP communication between the browser and the fake server, confirming that the traffic was redirected after the routing tables were poisoned.

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000	c2:02:37:fd:00:10	c2:02:37:fd:00:10	LOOP	60	Reply
2	1.733278	172.16.0.2	224.0.0.9	RIPv2	66	Response
3	3.737119	172.16.0.2	224.0.0.9	RIPv2	66	Response
4	5.748474	172.16.0.2	224.0.0.9	RIPv2	66	Response
5	6.658340	192.168.1.2	10.0.0.2	TCP	74	54068 → 80 [SYN] Seq=0 Win=64240 Len=0 MSS=1460 SACK_PERM=1 TSval=1606159760 TSecr=
6	6.658522	10.0.0.2	192.168.1.2	TCP	74	80 → 54068 [SYN, ACK] Seq=0 Ack=1 Win=65160 Len=0 MSS=1460 SACK_PERM=1 TSval=160606
7	6.688746	192.168.1.2	10.0.0.2	TCP	66	54068 → 80 [ACK] Seq=1 Ack=1 Win=64256 Len=0 TSval=1606159785 TSecr=1606068852
8	6.698864	192.168.1.2	10.0.0.2	HTTP	476	GET / HTTP/1.1
9	6.699281	10.0.0.2	192.168.1.2	TCP	66	80 → 54068 [ACK] Seq=1 Ack=411 Win=64768 Len=0 TSval=1606068892 TSecr=1606159785
10	6.699222	10.0.0.2	192.168.1.2	HTTP	254	HTTP/1.1 304 Not Modified
11	6.739220	192.168.1.2	10.0.0.2	TCP	66	54068 → 80 [ACK] Seq=411 Ack=189 Win=64128 Len=0 TSval=1606159835 TSecr=1606068892
12	7.744338	172.16.0.2	224.0.0.9	RIPv2	66	Response
13	9.748365	172.16.0.2	224.0.0.9	RIPv2	66	Response
14	9.991713	c2:02:37:fd:00:10	c2:02:37:fd:00:10	LOOP	60	Reply
15	11.746292	02:42:b2:9b:6c:00	c2:02:37:fd:00:10	ARP	42	Who has 172.16.0.1? Tell 172.16.0.2
16	11.751331	172.16.0.2	224.0.0.9	RIPv2	66	Response
17	11.752374	c2:02:37:fd:00:10	02:42:b2:9b:6c:00	ARP	60	172.16.0.1 is at c2:02:37:fd:00:10
18	13.755114	172.16.0.2	224.0.0.9	RIPv2	66	Response
19	15.768195	172.16.0.2	224.0.0.9	RIPv2	66	Response
20	15.898924	172.16.0.1	224.0.0.9	RIPv2	126	Response
21	16.875795	192.168.1.2	10.0.0.2	TCP	66	[TCP Keep-Alive] 54068 → 80 [ACK] Seq=410 Ack=189 Win=64128 Len=0 TSval=1606169974
22	16.875934	10.0.0.2	192.168.1.2	TCP	66	[TCP Keep-Alive ACK] 80 → 54068 [ACK] Seq=189 Ack=411 Win=64768 Len=0 TSval=16060615
23	17.760482	172.16.0.2	224.0.0.9	RIPv2	66	Response
24	19.763415	172.16.0.2	224.0.0.9	RIPv2	66	Response
25	19.997215	c2:02:37:fd:00:10	c2:02:37:fd:00:10	LOOP	60	Reply

Figure 12: TCP and HTTP traffic redirected to the fake Web server

As a countermeasure, RIPv2 authentication can be enabled on the routers. With authentication, routers only accept RIP updates from trusted devices, preventing forged routing updates from being installed in the routing table.

2.6.3 Effectiveness of the Attack

The success of the attack depends on the relative location of the browser, legitimate Web server and fake Web server. Since the attacker is placed near the fake Web server, it can advertise a route that makes the fake server appear as a better path to the target network.

The prefix length advertised by the attacker is also important. A more specific prefix is preferred due to the longest prefix match rule. Therefore, if the attacker advertises a more specific route than the legitimate one, the routers may select the malicious route.

RIP auto-summary can reduce the effectiveness of the attack because it summarizes routes at classful network boundaries. When auto-summary is disabled, more specific routes can be advertised, making the attack more effective.

Figures 13 and 14 show repeated RIPv2 Response messages, confirming that the attacker continuously injected routing updates into the network.

No.	Time	Source	Destination	Protocol	Length	Info
5	4.711615	192.168.23.2	224.0.0.9	RIPv2	106	Response
8	13.173999	192.168.23.2	224.0.0.9	RIPv2	66	Response
13	34.374262	192.168.23.2	224.0.0.9	RIPv2	106	Response
16	42.556733	192.168.23.2	224.0.0.9	RIPv2	66	Response
23	63.616773	192.168.23.2	224.0.0.9	RIPv2	106	Response
24	68.429552	192.168.23.2	224.0.0.9	RIPv2	66	Response

Figure 13: Filtered RIPv2 packets sent to multicast address 224.0.0.9

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000	c2:02:37:fd:00:00	c2:02:37:fd:00:00	LOOP	60	Reply
2	0.268414	c2:03:38:22:00:00	c2:03:38:22:00:00	LOOP	60	Reply
3	3.282191	c2:02:37:fd:00:00	CDP/VTP/DTP/PAgP/UD...	CDP	349	Device ID: R2 Port ID: FastEthernet0/0
4	3.481026	c2:03:38:22:00:00	CDP/VTP/DTP/PAgP/UD...	CDP	349	Device ID: R3 Port ID: FastEthernet0/0
5	4.711615	192.168.23.2	224.0.0.9	RIPv2	106	Response
6	9.998962	c2:02:37:fd:00:00	c2:02:37:fd:00:00	LOOP	60	Reply
7	10.275370	c2:03:38:22:00:00	c2:03:38:22:00:00	LOOP	60	Reply
8	13.173999	192.168.23.3	224.0.0.9	RIPv2	66	Response
9	20.005237	c2:02:37:fd:00:00	c2:02:37:fd:00:00	LOOP	60	Reply
10	20.270487	c2:03:38:22:00:00	c2:03:38:22:00:00	LOOP	60	Reply
11	30.003637	c2:02:37:fd:00:00	c2:02:37:fd:00:00	LOOP	60	Reply
12	30.270107	c2:03:38:22:00:00	c2:03:38:22:00:00	LOOP	60	Reply
13	34.374262	192.168.23.2	224.0.0.9	RIPv2	106	Response
14	40.001557	c2:02:37:fd:00:00	c2:02:37:fd:00:00	LOOP	60	Reply
15	40.272070	c2:03:38:22:00:00	c2:03:38:22:00:00	LOOP	60	Reply
16	42.556733	192.168.23.3	224.0.0.9	RIPv2	66	Response
17	49.995916	c2:02:37:fd:00:00	c2:02:37:fd:00:00	LOOP	60	Reply
18	50.267336	c2:03:38:22:00:00	c2:03:38:22:00:00	LOOP	60	Reply
19	60.000765	c2:02:37:fd:00:00	c2:02:37:fd:00:00	LOOP	60	Reply
20	60.272327	c2:03:38:22:00:00	c2:03:38:22:00:00	LOOP	60	Reply
21	63.284522	c2:02:37:fd:00:00	CDP/VTP/DTP/PAgP/UD...	CDP	349	Device ID: R2 Port ID: FastEthernet0/0
22	63.475356	c2:03:38:22:00:00	CDP/VTP/DTP/PAgP/UD...	CDP	349	Device ID: R3 Port ID: FastEthernet0/0
23	63.616773	192.168.23.2	224.0.0.9	RIPv2	106	Response
24	68.429552	192.168.23.3	224.0.0.9	RIPv2	66	Response
25	69.998044	c2:02:37:fd:00:00	c2:02:37:fd:00:00	LOOP	60	Reply
26	70.272005	c2:03:38:22:00:00	c2:03:38:22:00:00	LOOP	60	Reply
27	79.995333	c2:02:37:fd:00:00	c2:02:37:fd:00:00	LOOP	60	Reply
28	80.267995	c2:03:38:22:00:00	c2:03:38:22:00:00	LOOP	60	Reply

Figure 14: Repeated RIPv2 responses observed during the attack

2.6.4 AI Prediction vs Experimental Results

The AI predicted that the attacker would be more successful when advertising a more specific prefix, because routers prefer the longest prefix match. It also predicted that disabling RIP auto-summary would make the attack more effective, since the malicious route would not be summarized.

This prediction matched the experimental behavior. The captured RIP packets show that the attacker advertised specific routing information, and the browser traffic was redirected to the fake Web server. Therefore, the observed results confirm that prefix length and auto-summary directly influence route selection.

2.6.5 Feasibility

This attack is feasible in networks using RIP without authentication. Since RIP accepts routing updates from neighboring devices, an attacker can inject forged routes and redirect traffic.

However, in modern networks, RIP is less common and authentication can prevent this attack. Therefore, the attack is mainly feasible in poorly configured or legacy networks.

3 Firewalls

3.1 Classical firewalls versus Zone-Based Policy Firewalls

3.1.1 Firewall Validation

The firewall must enforce the following rules:

1. Block all traffic from INSIDE to OUTSIDE;
2. Allow only ICMP and HTTP traffic from OUTSIDE to the server;
3. Deny access to the firewall interfaces.

3.1.2 Classical Firewall (CBAC)

In the classical firewall, traffic filtering is based on ACLs combined with stateful inspection rules. As shown in the slides, inspection rules only track sessions and do not impose restrictions by themselves.

Initially, the configuration allowed access to multiple server ports because the ACL was too permissive. This confirms that **inspect rules do not have an implicit deny**, meaning that unwanted traffic may still pass if not explicitly blocked.

After correcting the ACLs, the observed behavior was:

1. Only HTTP and ICMP traffic to the server was allowed;
2. All other ports were classified as filtered by NMAP;
3. No traffic from INSIDE to OUTSIDE was possible due to explicit deny rules.

3.1.3 ZBPF

In ZBPF, interfaces are assigned to zones and policies are applied between zones. By default, traffic between different zones is dropped unless a policy is defined.

After configuration, the firewall behavior was:

1. Only explicitly permitted traffic (HTTP and ICMP to the server) was allowed;
2. All other inter-zone traffic was blocked by default;
3. Access to firewall interfaces was denied (controlled via the self zone).

Therefore, both firewalls satisfy the specification, but ZBPF enforces it more naturally due to its default deny model.

3.2 NMAP Operation

NMAP is used to discover hosts and scan ports by actively probing the network.

3.2.1 Host Discovery

When using `nmap -sn`, NMAP determines if hosts are active by sending:

1. ICMP Echo requests;
2. TCP SYN packets to common ports (e.g., port 80);
3. ARP requests in local networks.

A host is considered active if it replies to any of these probes.

3.2.2 Port Scanning

For port scanning, NMAP typically uses a TCP SYN scan:

1. SYN-ACK response $\Rightarrow$ port is open;
2. RST response $\Rightarrow$ port is closed;
3. No response or ICMP error $\Rightarrow$ port is filtered.

This method is efficient because it does not complete the TCP connection, allowing fast scanning of multiple ports.

3.3 CBAC vs ZBPF

The main differences between CBAC and ZBPF are:

1. **Configuration model:** CBAC is interface-based and relies on ACLs plus inspection rules, while ZBPF is zone-based and uses zone-pairs, class-maps, and policy-maps.
2. **Default behavior:** In CBAC, traffic may be allowed unless explicitly denied by ACLs. In ZBPF, traffic between zones is denied by default unless a policy allows it.
3. **Stateful inspection:** Both support stateful inspection, but in CBAC it depends on inspect rules applied to interfaces, whereas in ZBPF it is integrated into policy actions (inspect, pass, drop).
4. **Security and control:** ZBPF provides stricter and more modular control, since all inter-zone traffic must be explicitly defined. CBAC is less structured and more prone to configuration errors.

In conclusion, although both approaches can enforce the required policy, ZBPF offers a more robust, scalable, and secure solution.

3.4 Protecting a campus network using a ZBPF

3.5 Defense against DoS attacks

4 AAA

4.1 AAA with TACACS+

4.2 AAA with Radius

4.3 802.1X Authentication

5 Conclusion

6 Appendix